An aerial photograph of a vast, snow-covered mountain range. The peaks are rugged and partially covered in dark, rocky patches. The valleys are filled with deep snow, and the overall scene is a high-contrast, monochromatic landscape in shades of white, grey, and blue.

S&DS 265 / 565  
Introductory Machine Learning

# Stochastic Gradient Descent

September 23

Yale

# Goings on

- Assn 1 due tonight at 11:59pm ET
- Assn 2 will be posted by same time
- Quiz 1 graded
- Questions?

# Outline for today

- Stochastic gradient descent
- Application to logistic regression
- Regularization
- Learning rate and scaling
- Jupyter notebook example

# Stochastic gradient descent

- Suppose that we want to fit a really big model, where the number of samples  $n$  and number of variables  $p$  are very large
- The classical algorithms in standard software packages will fail
- How can we train such models?

# Example

- We want to classify ads according to whether or not they will be clicked on by a user
- We have a very large collection of training data
- Ads are represented in terms of a sparse list of features

1 | 5 : 1.1789641e-01    39 : 6.0373064e-02    45 : 1.3163488e-01

- The dataset is too large to load into memory, and the number of features is also very large
- New data are continually arriving
- How can we efficiently train a classifier?

# Online learning

We will introduce a method that

- Reads in the data points one at a time
- Updates the model for each sample
- Exploits sparsity of the features
- Uses little memory, never reads in the entire dataset

# Stochastic gradient descent

Initialize all parameters to zero:  $\beta_j = 0, j = 1, \dots, d$ .

Read through the data one record at a time, and update the model.

- 1 Read data item  $x$
- 2 Make a prediction  $\hat{y}(x)$
- 3 Observe the true response/label  $y$
- 4 Update the parameters  $\beta$  so  $\hat{y}$  is closer to  $y$

# Stochastic gradient descent

To begin, suppose we are doing *linear regression*. We initialize all parameters to zero:  $\beta_j = 0, j = 1, \dots, d$ .

We read through the data one record at a time, and update the model.

- 1 Read data item  $x$
- 2 Make a prediction  $\hat{y}(x) = \sum_{j=1}^p \beta_j x_j$
- 3 Observe the true response/label  $y$
- 4 Update the parameters  $\beta$  so  $\hat{y}$  is closer to  $y$



# SGD idea

Here's the idea:

- For each parameter  $\beta_j$ , see what happens to the loss if that parameter is increased a little bit.
- If the loss goes down (up), then increase (decrease)  $\beta_j$  proportionately
- Do this simultaneously for all of the parameters
- Rinse and repeat

# SGD idea

Change  $\beta_j$  by a little bit:

$$\beta_j \rightarrow \beta_j + \delta$$

What happens to the squared error?

$$\begin{aligned}(y - \hat{y})^2 &\rightarrow (y - \hat{y} - \delta x_j)^2 \\ &\approx (y - \hat{y})^2 - 2(y - \hat{y})\delta x_j\end{aligned}$$

We then change the parameter accordingly:

$$\beta_j \rightarrow \beta_j + 2\delta(y - \hat{y})x_j$$

# SGD for general loss

Suppose  $L(y, \beta^T x)$  is the loss for an input  $(x, y)$ , e.g.,  $(y - \beta^T x)^2$

SGD update:

$$\beta_j \leftarrow \beta_j - \delta \frac{\partial L(y, \beta^T x)}{\partial \beta_j}$$

$$\boldsymbol{\beta} \leftarrow \boldsymbol{\beta} - \delta \nabla_{\boldsymbol{\beta}} L(y, \boldsymbol{\beta}^T x) \quad (\text{vector notation})$$

- $\delta$  is the *learning rate* or “step size”
- Needs to be chosen carefully, getting smaller over time

# Gradient descent for general loss

If  $L(\beta)$  is the loss function over subset of training set:


$$\begin{aligned}L(\beta + \delta v) &\approx L(\beta) + \delta v^T \nabla L(\beta) \\L(\beta - \delta \nabla L(\beta)) &\approx L(\beta) - \delta \|\nabla L(\beta)\|^2\end{aligned}$$

This is why gradient descent is going downhill — if  $\delta$  is small enough.

“Batch” gradient descent uses the entire training set in each step of gradient descent.

*Stochastic* gradient descent computes a quick approximation to this gradient, using only a single or a small “mini-batch” of data points

# Batch vs. stochastic gradient descent



<https://wikidocs.net/3413>

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \delta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \delta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose  $y = 1$  and probability  $p(x)$  is high?

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \delta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \delta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose  $y = 1$  and probability  $p(x)$  is high? *small change*
- Suppose  $y = 1$  and probability  $p(x)$  is small?

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \delta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \delta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose  $y = 1$  and probability  $p(x)$  is high? *small change*
- Suppose  $y = 1$  and probability  $p(x)$  is small? *big change*  $\uparrow$
- Suppose  $y = 0$  and probability  $p(x)$  is small?



# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \delta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \delta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose  $y = 1$  and probability  $p(x)$  is high? *small change*
- Suppose  $y = 1$  and probability  $p(x)$  is small? *big change*  $\uparrow$
- Suppose  $y = 0$  and probability  $p(x)$  is small? *small change*
- Suppose  $y = 0$  and probability  $p(x)$  is big?

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \delta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \delta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose  $y = 1$  and probability  $p(x)$  is high? *small change*
- Suppose  $y = 1$  and probability  $p(x)$  is small? *big change*  $\uparrow$
- Suppose  $y = 0$  and probability  $p(x)$  is small? *small change*
- Suppose  $y = 0$  and probability  $p(x)$  is big? *big change*  $\downarrow$

# SGD: choice of learning rate

A conservative choice of learning rate is

$$\delta_t = \frac{1}{t}$$

A more aggressive choice is

$$\delta_t = \frac{1}{\sqrt{t}}$$

In practice: Try learning rates  $C/\sqrt{t}$  for different choices of  $C$ , and monitor the error

$$\frac{1}{T} \sum_{t=1}^T (Y_t - \hat{Y}_t)^2$$

# Demo

Open the demo notebook `sgd.ipynb` and follow along...

# SGD: choice of learning rate

Learning rate should scale as

$$\delta_t = \frac{1}{\sqrt{t}}$$

Problem: Some of the updates may be on different scales.

# SGD: choice of learning rate

Learning rate should scale as

$$\delta_t = \frac{1}{\sqrt{t}}$$

Problem: Some of the updates may be on different scales.

Solution: Let  $g_{tj} = \frac{\partial L(y_t, \beta^T x_t)}{\partial \beta_j}$

Scale gradients to get update rule

$$\beta_j \leftarrow \beta_j - \delta \frac{g_{tj}}{\sqrt{\sum_{s=1}^t g_{sj}^2}}$$

# SGD: scaling issues

For a linear model, the SGD update is

$$\beta_j \longleftarrow \beta_j - C_t x_j$$

If  $x_j$  increases by a factor of two, the parameter  $\beta_j$  should decrease by a factor of two.

This update doesn't respect that scaling

## SGD: scaling issues

Usual solution is to “standardize” each variable — subtract out the mean and divide by the standard deviation

$$x_j \leftarrow \frac{x_j - \text{mean}(x_j)}{\sqrt{\text{var}(x_j)}}$$

But this involves “looking ahead” to compute the mean and variance, and destroys the online property of the algorithm



# SGD: scaling issues

Usual solution is to “standardize” each variable — subtract out the mean and divide by the standard deviation

$$x_j \leftarrow \frac{x_j - \text{mean}(x_j)}{\sqrt{\text{var}(x_j)}}$$

But this involves “looking ahead” to compute the mean and variance, and destroys the online property of the algorithm

Solution: The mean and variance can be updated in an online manner, in constant time, by storing auxiliary variables for each component  $j$ .

# SGD: Regularization

A “ridge” penalty  $\lambda \sum_{j=1}^p \beta_j^2$  is easily handled.

Gradient changes by an additive term  $2\lambda\beta_j$ . Update becomes

$$\begin{aligned}\beta_j &\longleftarrow \beta_j + \delta \{ (y - p(x))x_j - \lambda\beta_j \} \\ &= (1 - \delta\lambda)\beta_j + \delta(y - p(x))x_j\end{aligned}$$

$$\beta_j x_j \longleftarrow (1 - \delta\lambda)\beta_j x_j + \delta(y - p(x))x_j^2$$

Observe that this “does the right thing” whether  $\beta_j$  wants to be large positive or negative.

- *The penalty shrinks  $\beta_j$  toward zero*

# What did we learn today?

- Stochastic gradient descent is a simple algorithm that can be applied to large classification and regression problems
- A parameter is updated according to how much the loss changes when that parameter is changed by a little bit
- This is the “go to” algorithm for fitting large or complex machine learning models
- Choosing the learning rate is a little tricky